

CLAUSE.

Every decision, cited. Auditable AI claims adjudication — every conclusion grounded in a verbatim policy clause, machine-verified before it is shown.

clause-4vv4.onrender.com · github.com/Dhevenddra/Clause

UNAUDITABLE AI CAN'T TOUCH REGULATED DECISIONS

- Claims adjudication is high-volume, manual, and expensive — a hair-on-fire workflow at every insurer and TPA.
- LLMs are blocked from it for one reason: **a wrong or unexplainable denial is a compliance event.**
- "Trust me" prompting doesn't survive an audit. Regulators want the clause, not the vibe.

"Why was this claim denied?"
– every regulator, every dispute, every time

UNEXPLAINABLE

DECISION + VERIFIED-CITATION LEDGER

The screenshot displays the CLAIM interface with two main sections: POLICY and CLAIM.

POLICY: # HORIZON HOME & CONTENTS POLICY - HP-2214 (Standard Form). Section 1 - Coverage. 1.1 This policy covers direct physical loss or damage to the insured dwelling and its contents caused by fire, lightning, storm, flood ingress through the roof, burst internal water pipes, theft following forcible entry, and impact by vehicles or falling trees. 1.2 Contents are covered up to a combined limit of ₹8,00,000 per event. Any single item of jewellery, electronics, or art is covered up to ₹75,000 unless individually scheduled in Annexure B.

CLAIM: # CLAIM FORM - CLM-88347. Policyholder: S. Aravindan · Policy: HP-2214 · Date of loss: ongoing, first noticed January 2026 · Date notified: 06 July 2026. Over the past several months the wooden flooring and skirting in two bedrooms have progressively warped, blackened, and crumbled. A carpenter inspected the rooms in June and confirmed extensive damp rot and termite infestation spreading from the bathroom wall, likely developing over one to two years due to a slow seepage.

ADJUDICATE: A large red stamp reads **DENY**. Below it, a risk score of 95/100 is shown with a red progress bar. Text indicates 2/2 citations verified in source · 9.05s. The model used is mow7a0s · 2026-07-11T06:56:02Z.

CITATION LEDGER: 2 decision points · each quote checked against Exhibit A.

- Paste (or load) a **policy + claim** → one structured Gemma call.
- **APPROVE / FLAG / DENY** + risk score + plain-language rationale.
- Every decision point cites a **verbatim clause**, stamped ✓ VERIFIED IN SOURCE
- Click a citation → the exact clause lights up in the policy.
- Downloadable audit record: input hashes, model ID, per-citation verification.

ZERO HALLUCINATION BY CONSTRUCTION

Trust is enforced by code, not promised by prompt.

1 · EXTRACT

Gemma must quote the policy **character-for-character** for every decision point. One structured call — no chains, no agents.

2 · VERIFY

`validation.py`: normalized exact-substring match against the source, mapped back to a character span. No fuzzy matching, no embeddings — a quote either exists or it doesn't.

3 · STAMP

Pass → ✓ **verified + clause highlight**.
Fail → ✗ **rejected in the open**,
excluded from the decision basis.
Never silently dropped.

8/8

CITATIONS VERIFIED
ACROSS LIVE DEMO
SCENARIOS

11/11

VALIDATION EDGE CASES PASS —
PARAPHRASE & FABRICATION REJECTED

1

MODEL CALL PER
ADJUDICATION

OPEN-WEIGHT GEMMA, AMD-READY BY DESIGN

- **Gemma 4 31B IT** (open weights) is the only adjudication model — served via the **Fireworks AI** API on a dedicated deployment.
- Same app runs Gemma on an **AMD Instinct MI300X** with vLLM on ROCm — 192 GB HBM3 runs it on a single card.
- Switching is **one environment variable**. No code changes.
- Open weights + self-hosting = the on-prem, air-gapped deployment story regulated insurers actually require.

```
OPENAI_BASE_URL=https://api.fireworks.ai/inference/v1  
MODEL_ID=.../gemma-4-31b-it-nvfp4#.../clause-gemma
```

↕ one env var

```
OPENAI_BASE_URL=http://mi300x-pod:8000/v1  
MODEL_ID=google/gemma-4-31b-it # vLLM · ROCm
```


THE ADJUDICATION LAYER INSURERS CAN DEPLOY

- Insurers process **hundreds of millions of claims a year**; routine adjudication still costs tens of dollars per claim in expert review time.
- The blocker was never model quality — it's **auditability**. CLAUSE makes every AI decision defensible line-by-line.
- Same trust core generalizes: vendor contracts, loan covenants, benefits eligibility, compliance review — any decision that must cite its source.
- **Ask:** design partners in claims ops; pilot on a real closed-claims corpus.



`clause-4vv4.onrender.com`
`github.com/Dhevenddra/Clause`